


# Deep Autoencoders for Anomaly Detection



# Auto = “Self”

- Autoencoders utilize self-supervision, a form of unsupervised learning that aims to reconstruct input data
- Formally, an autoencoder approximates the identity function:

$$f(x) = x$$


# Why Learn a Trivial Function?

The identity function is trivial; why learn it?


By placing constraints on a neural network architecture, specifically:

$$Z < X$$

Where **X** is the raw feature space and **Z** is a latent representation space, this becomes a non-trivial optimization problem where the model needs to learn some mapping from X to Z that **preserves information about an input x while reducing its dimensionality**.



In this image,  $X$  is the number of input/output dimensions ( $X=10$ ) and  $Z$  is the number of activation units in the “code” or encoding ( $Z=3$ ). Note that  $Z < X$ .



# Neural Network as a Composite Function

- A feed-forward neural network is a **composite function**, that is, it is composed of a **series of functions**.
- Each layer is a function.
- Autencoders exploit this property.

# Decomposition of a Feed-Forward NN

$$h(x) = z$$

# Decomposition of a Feed-Forward NN

$$h(x) = z$$

$$g(z) = \hat{y}$$



## Decomposition of a Feed-Forward NN

$$h(x) = z$$

$$g(z) = \hat{y}$$

$$f(x) = g(h(x)) = \hat{y}$$

## Decomposition of an Autoencoder

$$enc(x|\theta_{enc}) = z$$

## Decomposition of an Autoencoder

$$enc(x|\theta_{enc}) = z$$

$$dec(z|\theta_{dec}) = \hat{x}$$

## Decomposition of an Autoencoder

$$enc(x|\theta_{enc}) = z$$

$$dec(z|\theta_{dec}) = \hat{x}$$


$$f(x|\theta) = dec(enc(x|\theta_{enc})|\theta_{dec}) = \hat{x}$$

# MNIST Example

784-dim  $X$

10-dim  $Z$

784-dim  $X$



The 10-dimensional encoding, or “embedding”, contains enough information for the decoder network to approximate the input.

# Relation to Supervised Learning

In supervised learning, we learn some approximation of the function

$$F(x) = y$$

Where **x is a set of features** and **y is some known label or value**. The approximate function, or model, is defined by a set of parameters  $\theta$ :

$$f(x|\theta) = \hat{y}$$

# Supervised Learning

$$f(x|\theta) = \hat{y}$$

# Supervised Learning

$$f(x|\theta) = \hat{y}$$

$$L(y, \hat{y})$$



# Supervised Learning

$$f(x|\theta) = \hat{y}$$

$$L(y, \hat{y})$$

$$L(y, f(x|\theta))$$

## Supervised Learning

$$f(x|\theta) = \hat{y}$$

$$L(y, \hat{y})$$

$$L(y, f(x|\theta))$$

$$\min_{\theta} L(y, f(x|\theta))$$

# Self-Supervised Learning

In self-supervised learning, we learn some approximation of the identity function

$$f(x) = x$$

Where **x is a set of features**. This model is also defined by a set of parameters  $\theta$ :

$$f(x|\theta) = \hat{x}$$

## Supervised Learning

$$f(x|\theta) = \hat{y}$$

$$L(y, \hat{y})$$

$$L(y, f(x|\theta))$$

$$\min_{\theta} L(y, f(x|\theta))$$

## Self-Supervised Learning

$$f(x|\theta) = \hat{x}$$

## Supervised Learning

$$f(x|\theta) = \hat{y}$$

$$L(y, \hat{y})$$

$$L(y, f(x|\theta))$$

$$\min_{\theta} L(y, f(x|\theta))$$

## Self-Supervised Learning

$$f(x|\theta) = \hat{x}$$

$$L(x, \hat{x})$$

## Supervised Learning

$$f(x|\theta) = \hat{y}$$

$$L(y, \hat{y})$$

$$L(y, f(x|\theta))$$

$$\min_{\theta} L(y, f(x|\theta))$$

## Self-Supervised Learning

$$f(x|\theta) = \hat{x}$$

$$L(x, \hat{x})$$

$$L(x, f(x|\theta))$$

## Supervised Learning

$$f(x|\theta) = \hat{y}$$

$$L(y, \hat{y})$$

$$L(y, f(x|\theta))$$

$$\min_{\theta} L(y, f(x|\theta))$$

## Self-Supervised Learning

$$f(x|\theta) = \hat{x}$$

$$L(x, \hat{x})$$

$$L(x, f(x|\theta))$$

$$\min_{\theta} L(x, f(x|\theta))$$

# Applications

- Feature extraction
- De-noising
- Anomaly Detection
- Image Searches
- Image generation
- Image colorization
- Pixel Interpolation (increasing resolution of images)
- Synthetic Data generation
- Many more



# Anomaly Detection

Learn some

$$p(x) \approx P(x)$$

Where  $P(x)$  is the true probability distribution of some source of data,  $p(x)$  is a learned model. When

$$p(x) \ll 1$$

(or  $p(x)$  is very close to 0) for some  $x$ ,  $x$  is considered anomalous

## Reconstruction Loss Relation to $p(x)$

$$L(x, \hat{x}) \propto \frac{1}{p(x)}$$

Intuitively, autoencoders are good at reconstructing “typical” examples; they struggle to reconstruct anomalous examples.

# Autoencoder for Anomaly Detection

- Train autoencoder on a training set to convergence
- Define distribution of reconstruction loss on test set
- Define some threshold based on test set loss distribution
- For example, mean loss + 2 standard deviations
- Take new example, pass through autoencoder, compute reconstruction loss
- If loss > threshold, raise alarm
- Alternatively, inspect top N new examples ranked by reconstruction loss

# Code Examples

- [Credit Card Fraud Detection](#)
- [Handwritten Letter Detection with E-MNIST](#)